

V0: naive

N=64

```
[BENCH] B=2 H=12 N=64 D=64 iters=50
[BENCH] total: 26.527 ms | avg: 0.531 ms/call
```

N=128

```
[BENCH] B=2 H=12 N=128 D=64 iters=50
[BENCH] total: 95.248 ms | avg: 1.905 ms/call
```

N=256

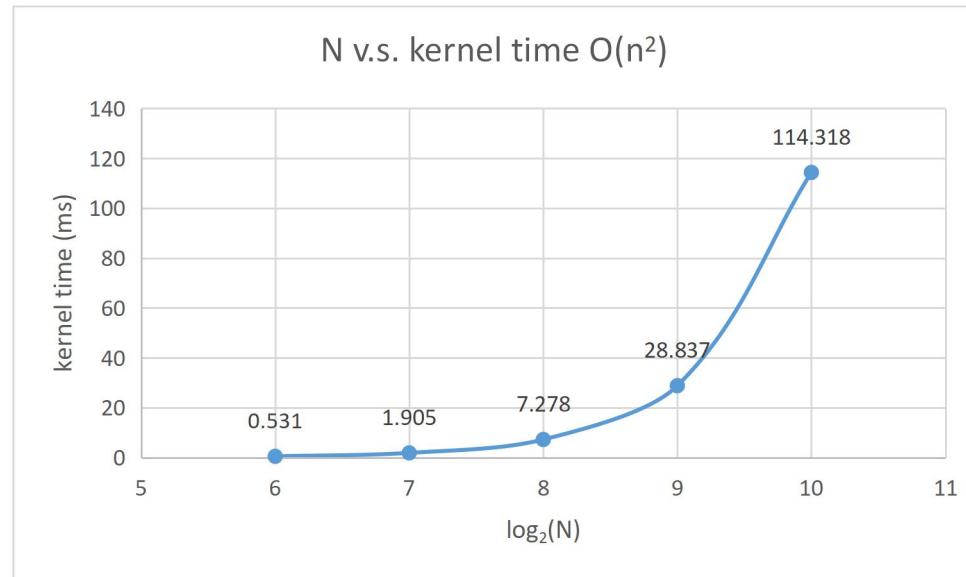
```
[BENCH] B=2 H=12 N=256 D=64 iters=50
[BENCH] total: 363.916 ms | avg: 7.278 ms/call
```

N=512

```
[BENCH] B=2 H=12 N=512 D=64 iters=50
[BENCH] total: 1441.836 ms | avg: 28.837 ms/call
```

N=1024

```
[BENCH] B=2 H=12 N=1024 D=64 iters=50
[BENCH] total: 5715.923 ms | avg: 114.318 ms/call
```

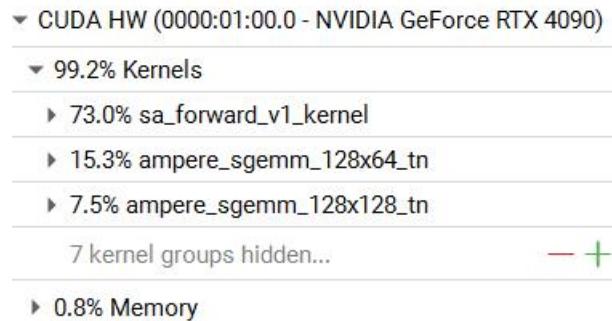


Bottleneck

▼ CUDA HW (0000:01:00.0 - NVIDIA GeForce RTX 4090)
▼ 99.9% Kernels
▶ 95.8% sa_forward_v0_kernel
▶ 2.3% ampere_sgemm_128x64_tn
▶ 1.2% ampere_sgemm_128x128_tn
7 kernel groups hidden... - +
▶ 0.1% Memory

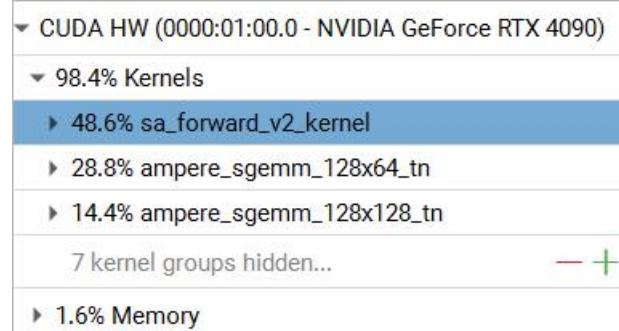
Inference loop: 8.423 sec for 10 batch

V1: online softmax , KV shared memory tile, single-pass output accumulation



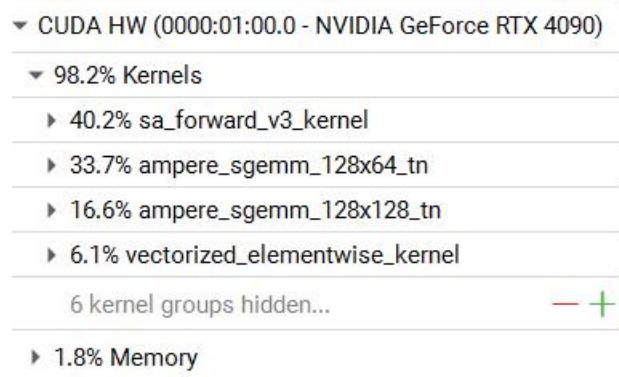
Inference loop: 1.276 sec for 10 batch

V2: warp-level design(one warp per query row), reuse K/V tile for multiple(8) Q rows, vectorized half2 loads and compute, warp shuffle reduction for dot-product



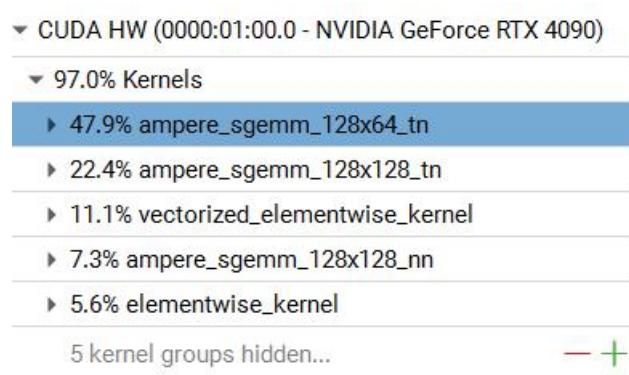
Inference loop: 686.089 ms for 10 batch

V3: Tensor Core calculate QxK dot products



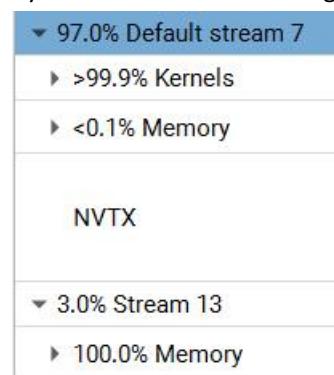
Inference loop: 578.692 ms for 10 batch

Pytorch: utilize Tensor Cores, highly optimized GEMM



Inference loop: 386.165 ms for 10 batch

Pytorch: multi-stream hiding memory latency



Inference loop: 365.383 ms for 10 batch